

PyCON HK

ORACLE®

Introduction of Graalpython to execute multiple languages with one processing system at high performance

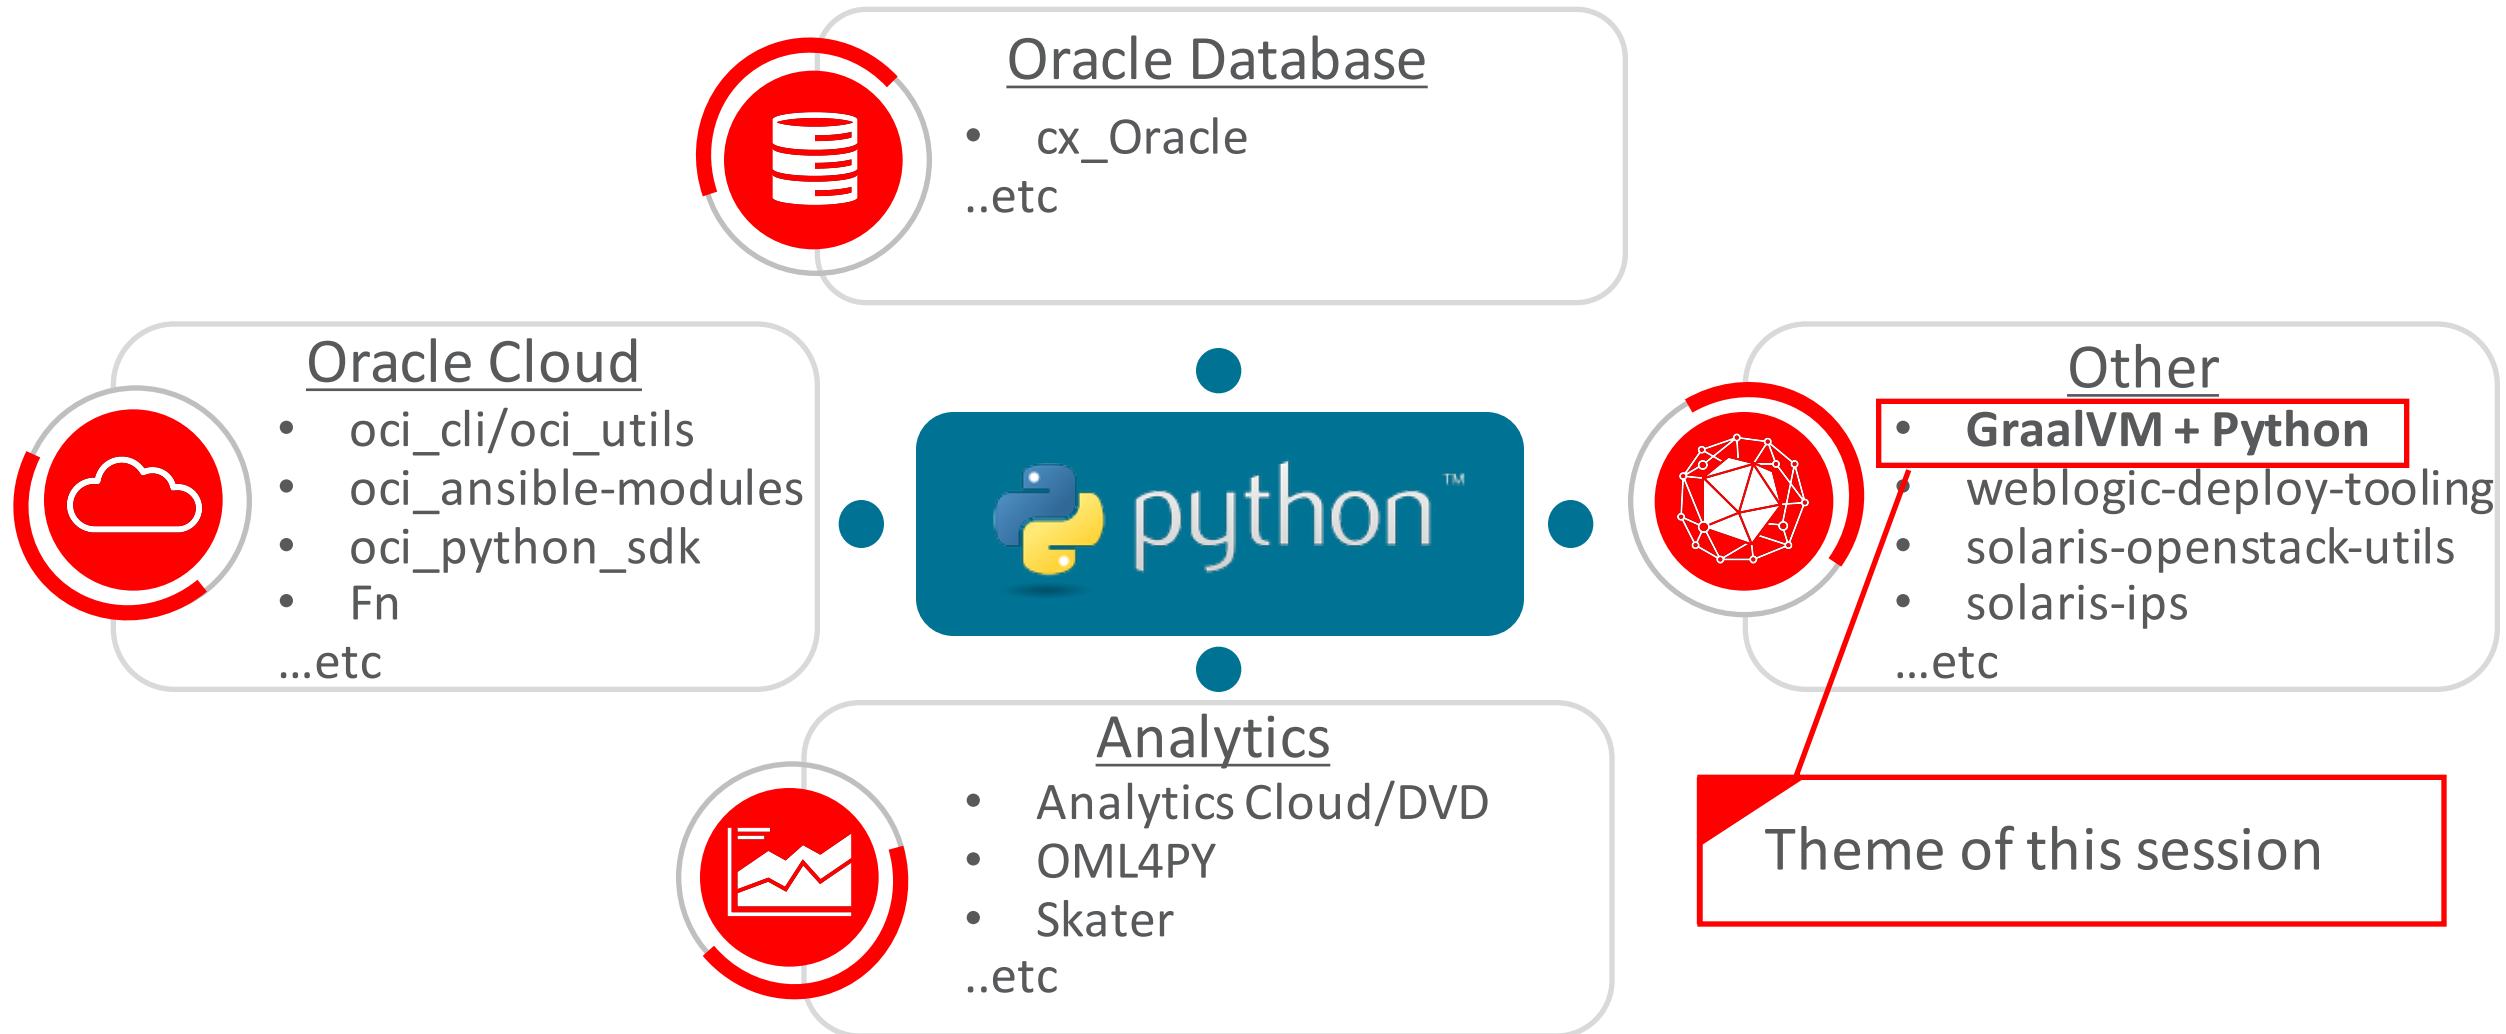
Amitpal Singh Dhillon

Regional Director, Asia-Pacific & Japan
Oracle Labs Research & Advanced Development

Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, timing, and pricing of any features or functionality described for Oracle's products may change and remains at the sole discretion of Oracle Corporation.

Relationship between Oracle and Python



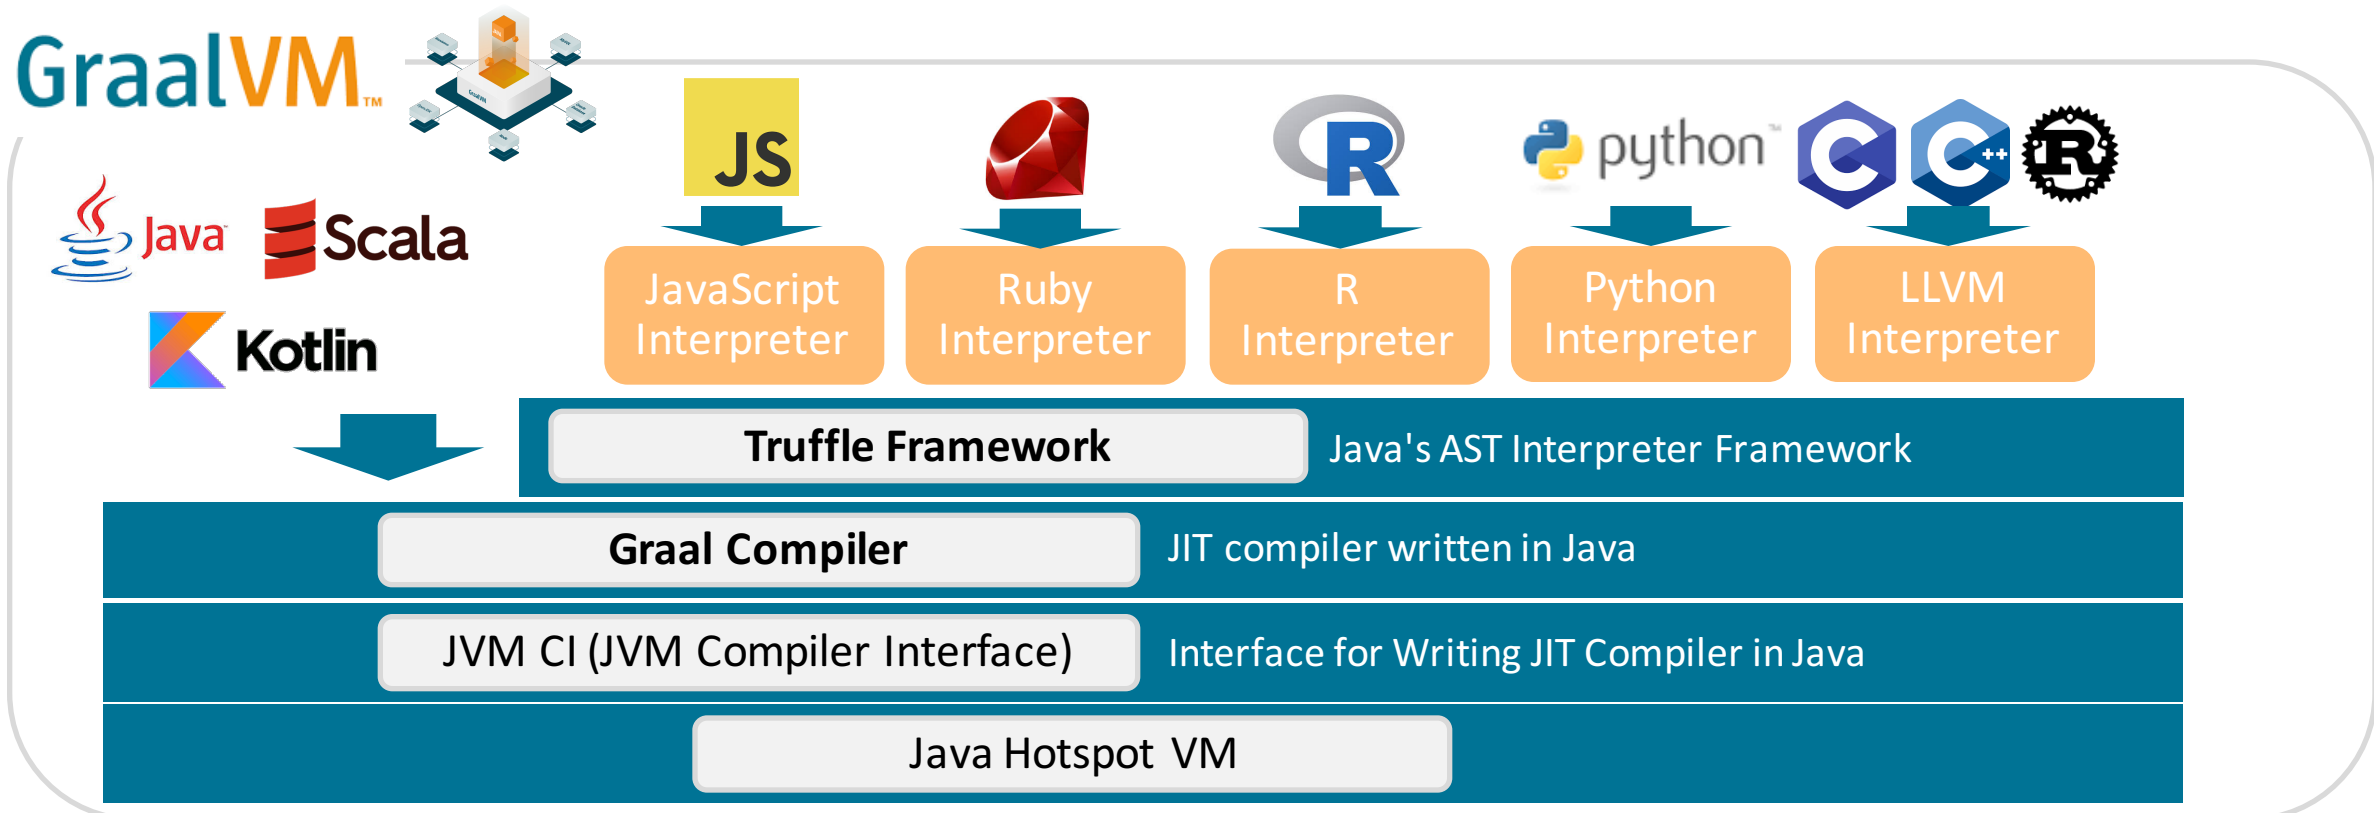
What is GraalVM

- A next generation VM that runs multiple languages fast with a single runtime



GraalVM Architecture


- Graal and Truffle for a High Performance Multilingual Executable VM Environment



GraalVM performance

- Good performance has been reported in various languages (Java / Scala / R / JavaScript)

Performance – Java and Scala




ORACLE

Copyright © 2014 Oracle and/or its affiliates. All rights reserved.

13

Performance – JavaScript



ORACLE

Copyright © 2014 Oracle and/or its affiliates. All rights reserved.

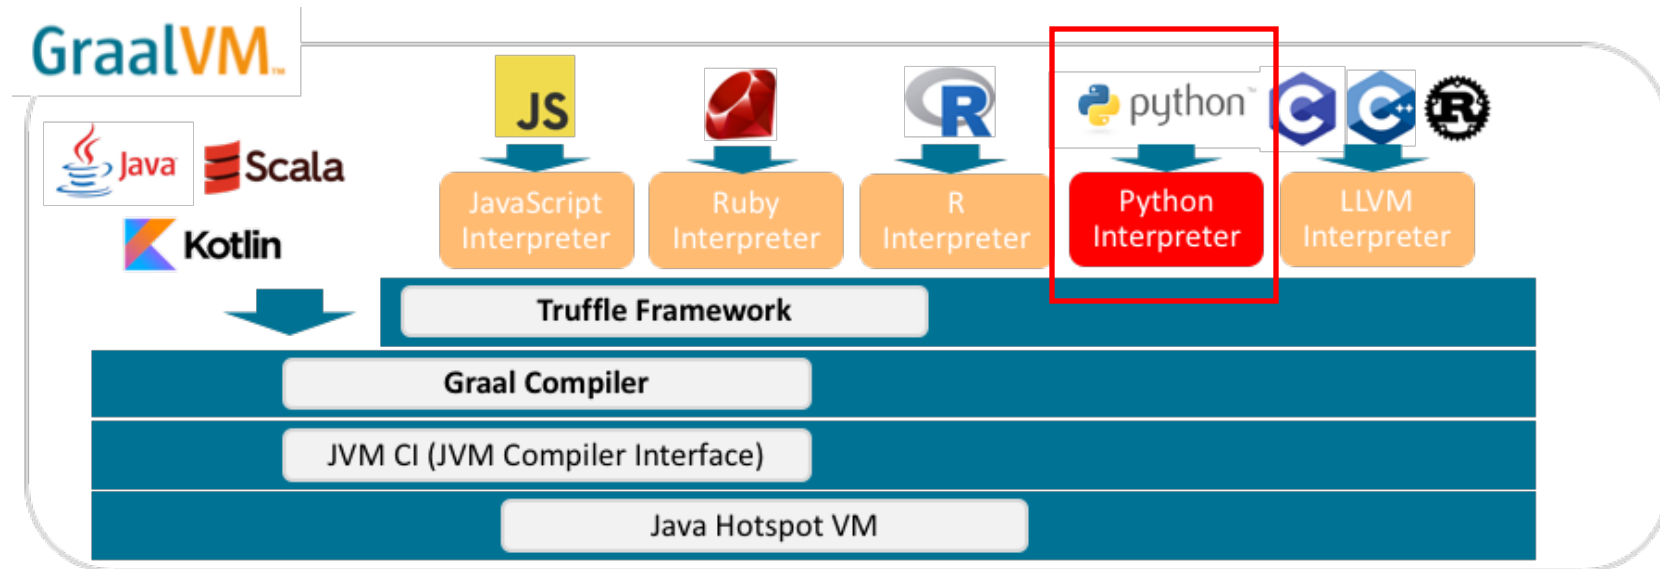
14

Ref. 「GraalVM: Multi-Language Execution」 Thomas Wuerthinger, Senior Research Director at Oracle
URL: https://www.slideshare.net/ThomasWuerthinger/jazoon2014-slides?from_action=save

ORACLE

What is graalpython

- Interpreter for Python 3 implementation built on GraalVM
- At the moment, “**early-stage experimental** implementation of Python”
- Ultimately, achieve full compatibility with Python 3.7 and aim for full support for Scipy and Scipy dependent libraries (Numpy, Pandas, Scikit-learn etc)



graalpython install & run

- Install with gu command of Graal and run GraalVM Python interpreter with graalpython command

Graalpython installation

```
$ gu -c install org.graalvm.python
```

```
$ graalpython --version
```

```
Graal Python 3.7.0 (GraalVM CE Native 1.0.0-rc15)
```

Run Graalpython

```
# Run a Python script
```

```
$ graalpython helloworld.py
```

```
# Run Python REPL
```

```
$ graalpython
```

```
Python 3.7.0 (Thu Apr 04 12:16:00 PDT 2019)
```

```
[GraalVM CE, Java 1.8.0_202] on darwin
```

```
Type "help", "copyright", "credits" or "license" for  
more information.
```

```
Please note: This Python implementation is  
in the very early stages, and can run little more than  
basic benchmarks at this point.
```

```
>>>
```


graalpython performance measurement

- In order to verify the high performance which is one of the attractiveness of graalpython, the following two benchmarks are executed on the Compute instance on OCI

① Some benchmark scripts run in the pyperformance (*) library

② Python script to perform pi calculation by Monte Carlo method

- Compare three Python implementations (Cpython, Jython, PyPy) with graalpython



Cpython 3.7.1



Jython 2.7



PyPy 3.6



Graal Python 3.7.0
(GraalVM CE Native 1.0.0-rc15)



Virtual
Machine

VM.Standard2.16

CPU	32 cores
メモリ	2
OS	Ubuntu 18.04 LTS

[Reference] Python Implementations

Implementation	Latest Version	Standard Python Compatible version	Feature
CPython	3.7.3	(Standard Python implementation)	Python language reference implementation https://www.python.org/
Jython	2.7.0	2.7	Java implementation of Python https://www.jython.org/
PyPy	3.6	3.6	A Python implementation of RPython designed to accelerate CPython https://pypy.org/index.html
graalpython	3.7 GraalVM 1.0.0-rc16	3.7 ✂ However, because it is experimental, it is not yet fully compatible	Python implementation that runs on Graal VM https://github.com/graalvm/graalpython
Zippy	3.7	3.7	Python implementation that works on Truffle Framework https://github.com/seuresystemslab/zippy

graalpython performance measurement – pyperformance

The following three scripts in pyperformance are slightly modified and executed because there are limitations on libraries that can be operated with graalpython

bm_spectral_norm.py

Script to solve "Hundred-Dollar, Hundred-Digit Challenge Problems"

<https://pyperformance.readthedocs.io/benchmarks.html#spectral-norm>

bm_nqueens.py

Script that solves N-Queen problem with simple solver


<https://pyperformance.readthedocs.io/benchmarks.html#nqueens>

bm_fannkuch.py

Game script for computer language benchmark

<https://pyperformance.readthedocs.io/benchmarks.html#fannkuch>

Results
pyperformance benchmark tests



pyperformance tests

- graalpython outperforms PyPy but gives better performance than Cpython
- "bm_nqueens.py" uses Generator extensively, and it seems that the execution speed of CPython, which is the standard implementation, is optimized and executed more than graalpython.

[Reference] GraalVM Debugging Tools - chrome-devtools

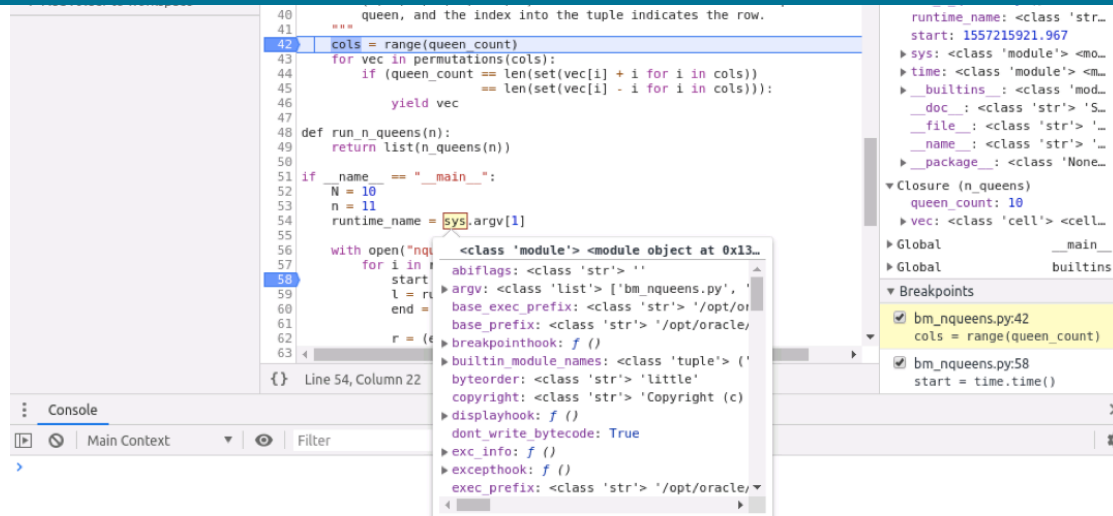
```
$ graalpython --jvm --inspect <your_python_script_name>.py
```

Debugger listening on port 9229.

To start debugging, open the following URL in Chrome:

`chrome-devtools://devtools/bundled/js_app.html?ws=127.0.0.1:9229/6537cf78-154605af4255`

- Settings Break Point
- Executing Python functions on browser console
- Referencing to variables



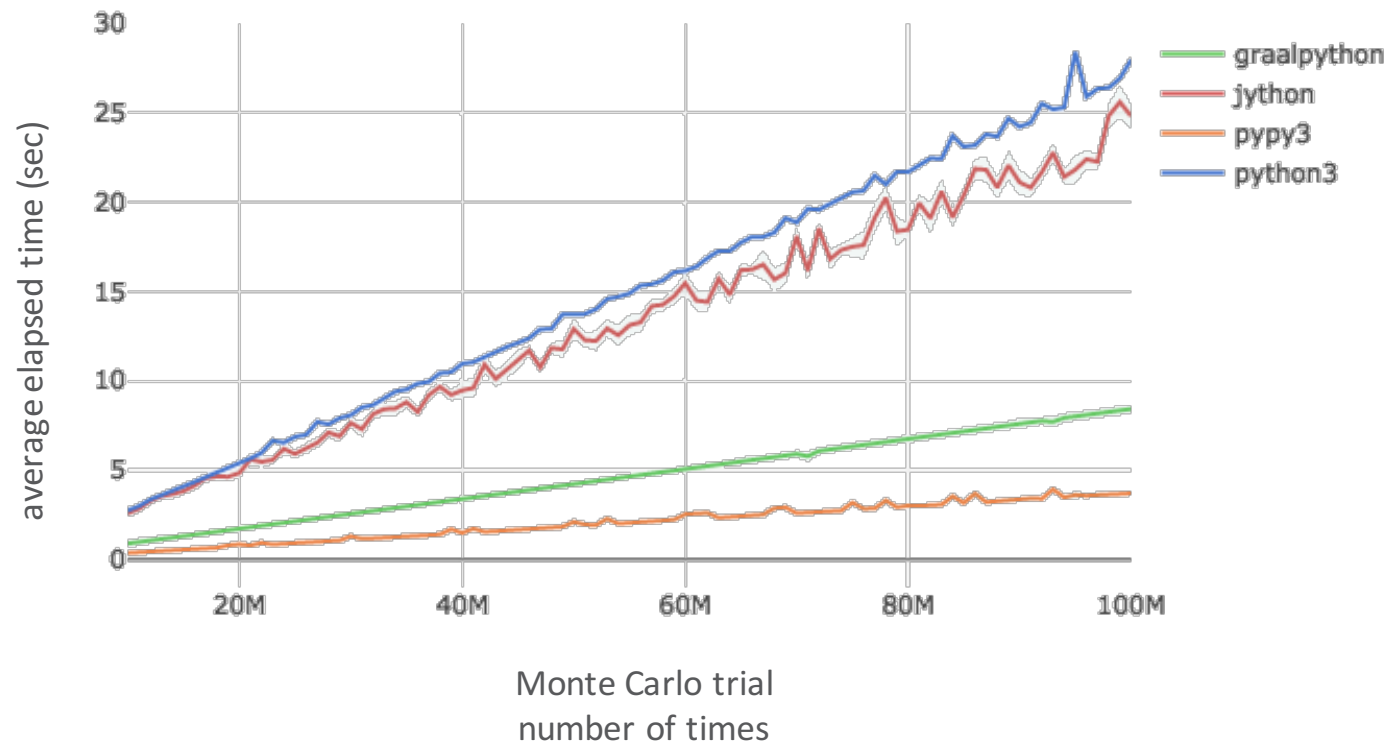
Profiles	Self Time	Total Time	Function
CPU PROFILES	20591.0 ms 76.50 %	20591.0 ms 76.50 %	n_queens.<locals>.<genexp>bm_nqueens.py:44
	2064.0 ms 7.67 %	2308.0 ms 8.57 %	permutations
	1515.0 ms 5.63 %	1515.0 ms 5.63 %	n_queens.<locals>.<genexp>bm_nqueens.py:45
	244.0 ms 0.91 %	244.0 ms 0.91 %	permutations.<locals>.<genexp>bm_nqueens.py:25
	72.0 ms 0.27 %	24486.0 ms 90.97 %	n_queens
	0 ms 0 %	24486.0 ms 90.97 %	<module 'bm_nqueens.py'>
	0 ms 0 %	24486.0 ms 90.97 %	<module 'bm_nqueens.py'>
	0 ms 0 %	24486.0 ms 90.97 %	<module 'bm_nqueens.py'>
	0 ms 0 %	24486.0 ms 90.97 %	<module 'bm_nqueens.py'>
	0 ms 0 %	24486.0 ms 90.97 %	<module 'bm_nqueens.py'>
	0 ms 0 %	24486.0 ms 90.97 %	<module 'bm_nqueens.py'>
	0 ms 0 %	24486.0 ms 90.97 %	run_n_queens
	0 ms 0 %	24486.0 ms 90.97 %	run_n_queens
	0 ms 0 %	24486.0 ms 90.97 %	run_n_queens
	0 ms 0 %	24486.0 ms 90.97 %	n_queens
	0 ms 0 %	22555.0 ms 83.79 %	n_queens
	0 ms 0 %	22555.0 ms 83.79 %	n_queens
	0 ms 0 %	1546.0 ms 5.74 %	n_queens

- Profiling Python scripts to see which functions have consumed how much time

graalpython performance measurement-Monte Carlo method 1

Execute script to approximate pi by Monte Carlo method (*)

Execution time of pi calculation by Monte Carlo method




- Even in the Monte Carlo calculation, although it is not comparable to PyPy, the result is higher than CPython and Jython.
- A simple CPU bound process without I / O such as Monte Carlo calculation may be considered as an optimized calculation.

(*) ref. **Monte Carlo method (Wikipedia)**
https://en.wikipedia.org/wiki/Monte_Carlo_method

graalpython performance measurement-Monte Carlo method 2

Execute script that approximates pi by multi-thread (*) by Monte Carlo method

Execution time of pi calculation by Monte Carlo method (multi-thread)




- Graalpython seems to be a processing system with GIL as well as CPython and PyPy, because the processing time hardly changes even if the number of threads is increased.
- On the other hand, Jython that does not have GIL has a shorter processing time as the number of threads increases, and it can be seen that multi-thread execution can be performed with multiple cores.

(*) Implementation using threading's `threading.Thread` class, which is a Python standard library
<https://docs.python.org/3/library/threading.html#threading.Thread>

graalpython performance measurement-Monte Carlo method 2

Execute script that approximates pi by multi-thread (*) by Monte Carlo method

Execution time of pi calculation by Monte Carlo method (multi-thread)



- Graalpython seems to be a processing system with GIL as well as CPython and PyPy, because the processing time hardly changes even if the number of threads is increased.
- On the other hand, Jython that does not have GIL has a shorter processing time as the number of threads increases, and it can be seen that multi-thread execution can be performed with multiple cores.

(*) Implementation using threading's `threading.Thread` class, which is a Python standard library
<https://docs.python.org/3/library/threading.html#threading.Thread>

graalpython and GIL (Global Interpreter Lock)

- GIL (Global Interpreter Lock) is a protection mechanism to limit the number of threads that can be executed at one time on the interpreter to one
 - It can simplify low-level processing such as memory management and cooperation of C extension
- Graalpython uses GIL in the C extension part (*), and multi-thread processing is executed by one CPU core as CPython

Implementation	With or without GIL
CPython	With GIL
Jython	Without GIL
PyPy	With GIL ※ However, Software Transactional Memory (STM) can avoid GIL http://doc.pypy.org/en/latest/stm.html
graalpython	With GIL

(*) Ref: <https://github.com/graalvm/graalpython/blob/master/graalpython/com.oracle.graal.python.cext/include/pystate.h#L378-L432>

Another attraction of graalpython (Polyglot) 1

- The attraction of graalpython is not only high performance, but also the ability to run Polyglot (multilingual execution) programs easily and without any loss of performance

※ However, as of May 2019 graalpython does not support Scikit-Learn

[Example]
data analysis with R and Python



tidyverse library Modern
data shaping used



Machine Learning Model
Execution with
Scikit-Learn Library

sample_polyglot.py

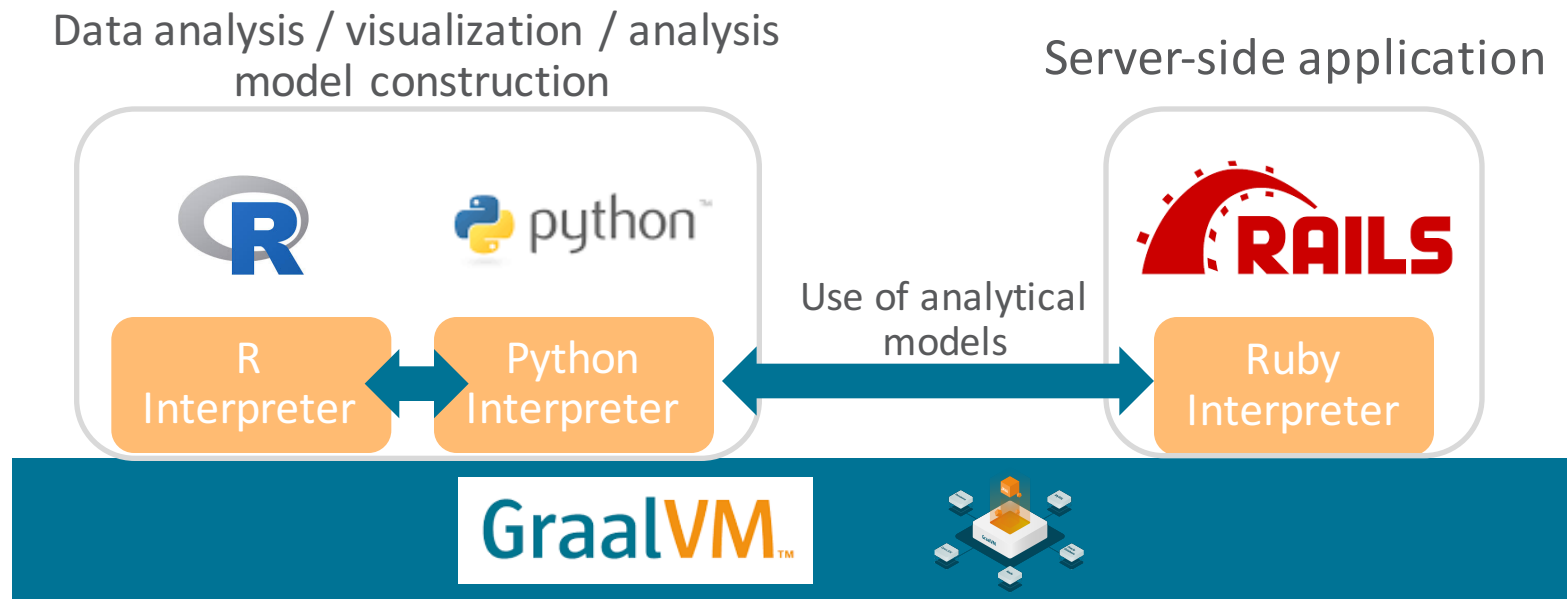
```
import polyglot
import sklearn
```

```
df = polyglot.eval(string="
    library(tidyverse)
    iris %>% dplyr::group_neset(Species) %>% dplyr::summarise_all(mean) ...
")
...
```

```
clt = sklearn.ensemble.RandomForestClassifier()
clt.fit(df["X"], df["y"])
clf.feature_importances_
```

Another attraction of graalpython (Polyglot) 2

- Furthermore, combining with JavaScript and Ruby makes server-side collaboration easy and reduces the restriction on programming language selection for creating applications
- For example, the following machine learning model embedded application can be configured in the future



Summary

- graalpython is currently an experimental implementation, but performs better than CPython (the standard Python implementation)
- Not only performance but also Polyglot is one of the attractiveness of graalpython
 - In the area of Analytics, which is the main usage scene of Python, it is also possible to analyze combining R and Python, for example
 - It is possible to use different languages according to the characteristics of the program on a single Graal VM



Try graalpython ! GraalVM™ + python™